Journal of Statistical Software

MMMMMM YYYY, Volume VV, Issue II. doi: 10.18637/jss.v000.i00

peppwR: Simulation-Based Power Analysis for Phosphoproteomics Experiments in R

Dan MacLean
The Sainsbury Laboratory

Abstract

Phosphoproteomics experiments require careful experimental design to ensure adequate statistical power, yet no dedicated tools exist for power analysis that account for the unique statistical properties of peptide-level quantification data. We present **peppwR**, an R package for simulation-based power analysis tailored to phosphoproteomics. **peppwR** fits distributions to pilot data on a per-peptide basis, capturing the heterogeneity in abundance and variance across the peptidome. It supports multiple statistical tests (Wilcoxon, bootstrap-t, Bayes factor), tracks missingness patterns common in mass spectrometry data, and provides FDR-aware power calculations. Users can determine required sample sizes, estimate power for planned experiments, or identify minimum detectable effect sizes.

Keywords: power analysis, phosphoproteomics, mass spectrometry, experimental design, R.

1. Introduction

Statistical power—the probability of correctly rejecting a false null hypothesis—represents one of the most fundamental concepts in experimental design (**?**). Introduced by **?** and formalized by **?**, power analysis provides the theoretical framework for determining appropriate sample sizes, estimating the likelihood of detecting true effects, and identifying minimum detectable effect sizes. The consequences of inadequate power are profound: under-powered studies waste resources and risk missing genuine biological discoveries, while over-powered designs are economically inefficient and may detect biologically trivial effects (**??**).

Despite widespread adoption in clinical trials (**?**) and genomics (**?**), power analysis remains systematically neglected in proteomics research. Recent surveys indicate that fewer than 15% of published proteomics studies report formal power calculations (**?**), compared to over 80% in clinical research (**?**). This disparity persists despite the field's evolution toward increasingly quantitative approaches and the availability of sophisticated statistical methods for proteomics data analysis (**?**).

Mass spectrometry-based proteomics presents unique statistical challenges that violate fundamental assumptions of traditional power analysis. First, peptide abundances exhibit extreme heterogeneity, with concentrations spanning 6-8 orders of magnitude within a single experiment (**??**). This heterogeneity extends beyond mean abundances to encompass dramatic differences in variance structure across the peptidome (**??**). Second, abundance distributions deviate substantially from normality, typically exhibiting right skewness better characterized by gamma, lognormal, or inverse Gaussian distributions (**??**). These distributional properties reflect both biological variation and the multiplicative nature of mass spectrometry quantification processes (**?**).

Third, missing values in proteomics follow systematic, non-random patterns that fundamentally alter statistical inference. Low-abundance peptides are disproportionately affected by detection limits, creating missing-not-at-random (MNAR) data structures (**??**). The proportion of missing values can exceed 50% in discovery proteomics experiments, with missingness rates negatively correlated with abundance levels—a pattern that violates the missing-completely-at-random (MCAR) assumptions underlying most statistical tests (**??**). Fourth, phosphoproteomics experiments routinely quantify thousands of peptides simultaneously, necessitating multiple testing corrections that substantially reduce effective statistical power (**??**).

Current power analysis tools address these proteomics-specific challenges incompletely or not at all (Table **??**). **Perseus** (**?**), the dominant analysis platform in proteomics with over 10,000 citations, provides comprehensive statistical analysis including t-tests, ANOVA, and volcano plots, but offers no power analysis functionality. Users must rely on generic tools or ad-hoc approaches for experimental design decisions. The **clippda** package (**?**) represents the only proteomics-focused power analysis tool, but was specifically designed for SELDI-TOF peak data with strong assumptions of homogeneous variance across features—assumptions that poorly reflect modern LC-MS/MS data characteristics.

G*Power (**?**), while widely used across scientific disciplines with over 30,000 citations, assumes normality and homogeneous variance, making it inappropriate for proteomics applications. Its power calculations for t-tests and ANOVA fundamentally misestimate required sample sizes when applied to non-normal, heterogeneous proteomics data. **MultiPower** (**?**) addresses power analysis in multi-omics contexts but requires simultaneous data from multiple platforms (genomics, transcriptomics, metabolomics), limiting its applicability to proteomics-only experiments. **PowerAtlas** (**?**) focuses on pathway-level analysis rather than individual protein or peptide quantification.

The absence of appropriate tools has led researchers to either skip power analysis entirely or apply inappropriate methods, potentially contributing to the reproducibility challenges observed in proteomics research (**??**). This methodological gap is particularly problematic given the high costs associated with proteomics experiments and the limited sample availability in many clinical contexts (**?**).

We present **peppwR**, an R package that directly addresses these limitations through simulation-based power analysis tailored specifically to mass spectrometry proteomics data. **peppwR** implements per-peptide distribution fitting to capture abundance heterogeneity, supports multiple statistical tests appropriate for non-normal data, incorporates realistic missing data patterns, and provides FDR-aware power calculations for high-throughput experiments. The package operates in two complementary modes: aggregate analysis for preliminary design decisions without pilot data, and per-peptide analysis that leverages pilot data to provide detailed, peptide-specific power estimates.

peppwR answers the three fundamental experimental design questions: "What sample size do I need to achieve target power?", "What statistical power will my planned experiment achieve?", and "What is the minimum detectable effect size given my sample size constraints?" By addressing the unique statistical properties of proteomics data, **peppwR** enables more informed experimental design decisions and contributes to improving the reproducibility and efficiency of proteomics research.

Table 1: Comparison of power analysis tools for proteomics applications.

Tool	Proteomics Focus	Per-feature Modeling	Non-normal Distributions	Missing Data Handling	Multiple Testing	Statistical Tests
peppwR	✓	✓	✓	✓	✓	Wilcoxon, Bootstrap-t Bayes Factor
Perseus	✓	✗	✓	✓	✓	t-test, ANOVA Volcano plots
clippda	✗[1]	✗	✗	✗	✗	t-test
MultiPower	✗[2]	✓	✗	✗	✓	Various
G*Power	✗	✗	✗	✗	✗	t-test, ANOVA Correlation

2. Methods and Implementation

peppwR implements simulation-based power analysis through a modular framework that addresses the unique statistical challenges of phosphoproteomics data. The package operates in two complementary modes: aggregate analysis for preliminary design decisions without pilot data, and per-peptide analysis that leverages pilot data to provide detailed, peptide-specific power estimates across the heterogeneous peptidome.

2.1. Statistical Framework

Statistical power, defined as $P(\text{reject } H_0|H_1 \text{ true}) = 1-\beta$ (where β is the Type II error rate), provides the theoretical foundation for experimental design decisions. For a given statistical test with significance level α, power depends on the effect size δ, sample size n, and the underlying population distributions. In the context of proteomics experiments comparing two groups, we test the null hypothesis $H_0 : \mu_{\text{treatment}} = \mu_{\text{control}}$ against the alternative $H_1 : \mu_{\text{treatment}} \neq \mu_{\text{control}}$.

peppwR accommodates three fundamental experimental design questions through constrained optimization. When `find = "sample_size"`, the package determines the minimum n required to achieve target power π_{target}:

$$n^* = \min\{n : P(\text{reject } H_0|\delta, n) \geq \pi_{\text{target}}\} \tag{1}$$

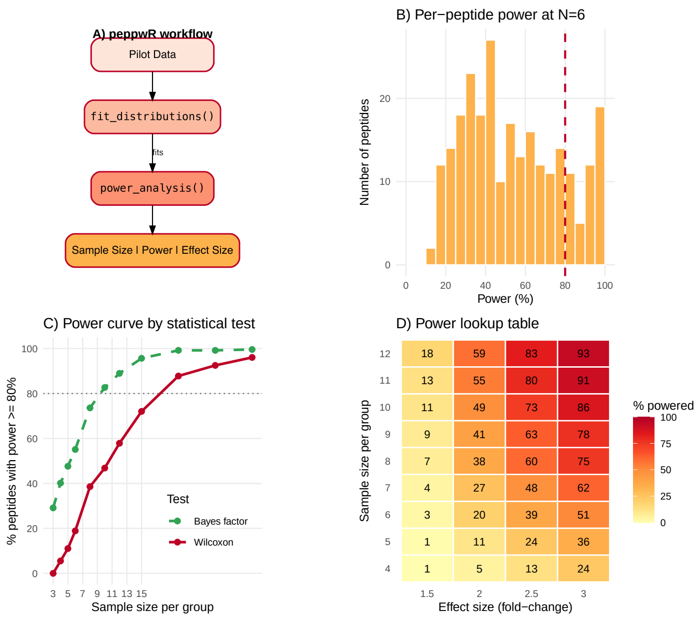

Figure 1: Overview of the **peppwR** workflow showing the two-stage process: distribution fitting followed by power analysis via Monte Carlo simulation.

When `find = "power"`, it estimates the expected power for a fixed sample size:

$$\pi(n, \delta) = P(\text{reject } H_0 | \delta, n) \tag{2}$$

When `find = "effect_size"`, it identifies the minimum detectable effect size:

$$\delta^* = \min\{\delta : P(\text{reject } H_0 | \delta, n) \geq \pi_{\text{target}}\} \tag{3}$$

The effect size is parameterized as a multiplicative factor δ such that treatment group abundances are scaled by δ relative to control abundances. This multiplicative parameterization naturally accommodates the log-scale nature of mass spectrometry quantification, where biological effects typically manifest as fold-changes rather than additive differences.

In per-peptide mode, power calculations are performed independently for each peptide i with fitted distribution parameters θ_i, yielding peptide-specific power estimates $\pi_i(n, \delta)$. The package then reports the proportion of peptides achieving target power:

$$P_{\text{peptidome}}(n, \delta, \pi_{\text{target}}) = \frac{1}{K} \sum_{i=1}^{K} \mathbf{1}[\pi_i(n, \delta) \geq \pi_{\text{target}}] \tag{4}$$

where K is the total number of peptides and $\mathbf{1}[\cdot]$ is the indicator function.

For FDR-aware analysis, power is evaluated after Benjamini-Hochberg correction. Let $p_{(1)} \leq p_{(2)} \leq \ldots \leq p_{(m)}$ denote the ordered p-values from m tests. The FDR-controlling procedure rejects all hypotheses $H_{(i)}$ for $i \leq k^*$, where:

$$k^* = \max\left\{i : p_{(i)} \leq \frac{i \cdot \alpha}{m}\right\} \tag{5}$$

Power under FDR correction becomes the expected proportion of true positives among discoveries, accounting for the adaptive threshold imposed by the multiple testing procedure.

2.2. Distribution Fitting

peppwR models peptide abundance distributions using maximum likelihood estimation across multiple parametric families. The core distributions supported are gamma (Gamma(α, β)), lognormal (LogNormal(μ, σ^2)), normal (Normal(μ, σ^2)), and inverse Gaussian (InverseGaussian(μ, λ)). These distributions capture the right-skewed, positive-valued nature of mass spectrometry intensities while accommodating varying degrees of skewness and heteroscedasticity observed across the peptidome.

For each peptide i with observed abundances $\mathbf{x}_i = \{x_{i1}, x_{i2}, \ldots, x_{in_i}\}$, the package fits each candidate distribution with probability density function $f_k(x|\theta_k)$ by maximizing the log-likelihood:

$$\hat{\theta}_{k,i} = \arg\max_{\theta_k} \sum_{j=1}^{n_i} \log f_k(x_{ij}|\theta_k) \tag{6}$$

Model selection employs the Akaike Information Criterion (AIC), selecting the distribution with minimum AIC score:

$$\hat{k}_i = \arg\min_{k} \left\{ -2\sum_{j=1}^{n_i} \log f_k(x_{ij}|\hat{\theta}_{k,i}) + 2p_k \right\} \tag{7}$$

where p_k is the number of parameters in distribution k.

The gamma distribution, parameterized by shape $\alpha > 0$ and rate $\beta > 0$, provides flexibility in modeling right-skewed abundance data with controlled variance-to-mean relationships. The lognormal distribution naturally accommodates multiplicative biological processes underlying protein expression, while the inverse Gaussian distribution offers an intermediate option between gamma and lognormal with distinct tail behavior. Additional candidate distributions including Weibull, exponential, and beta (for transformed data) are available through the **univariateML** interface.

Parameter estimation leverages the **fitdistrplus** package for standard distributions and **univariateML** for specialized cases. Both packages implement numerically stable algorithms with appropriate handling of boundary conditions and convergence diagnostics. When maximum likelihood estimation fails due to insufficient data, extreme parameter estimates, or numerical instability, **peppwR** provides three fallback strategies controlled by the `on_fit_failure` argument.

The `"exclude"` strategy removes problematic peptides from analysis, reporting the exclusion count for transparency. The `"empirical"` strategy substitutes parametric distributions with empirical bootstrap resampling from observed values, preserving the exact empirical distribution at the cost of reduced statistical efficiency. The `"lognormal"` strategy fits a moment-matched lognormal distribution using method-of-moments estimates, providing a robust fallback when maximum likelihood fails but ensuring all peptides remain in analysis.

Distribution fitting diagnostics include goodness-of-fit statistics (Kolmogorov-Smirnov, Anderson-Darling), QQ plots comparing theoretical and empirical quantiles, and density overlay plots visualizing the fitted distribution against observed data histograms. These diagnostics enable users to assess the adequacy of distributional assumptions and identify peptides where parametric modeling may be inappropriate.

2.3. Simulation Engine

Power estimation employs Monte Carlo simulation to accommodate the complex statistical properties of proteomics data that preclude analytical power calculations. The simulation framework generates synthetic datasets under specified experimental conditions, applies statistical tests, and estimates power as the proportion of simulations achieving statistical significance.

For each peptide i with fitted distribution $f_i(\cdot|\hat{\theta}_i)$ and each candidate sample size n, the simulation algorithm proceeds as follows:

1. Generate n control group samples: $X_{i,\text{control}}^{(s)} \sim f_i(\cdot|\hat{\theta}_i)$ for $s = 1, \ldots, n$
2. Generate n treatment group samples: $X_{i,\text{treatment}}^{(s)} \sim f_i(\cdot|\delta \cdot \hat{\theta}_i)$ for $s = 1, \ldots, n$, where δ represents the multiplicative effect size
3. Apply missing data patterns if `include_missingness = TRUE`, randomly removing observations according to the peptide-specific missingness rate r_i
4. Compute the test statistic and p-value using the specified statistical test
5. Record rejection outcome: $R^{(j)} = \mathbf{1}[p^{(j)} \leq \alpha]$ for simulation j
6. Repeat for n_{sim} independent simulations

Power is estimated as:

$$\hat{\pi}_i(n, \delta) = \frac{1}{n_{\text{sim}}} \sum_{j=1}^{n_{\text{sim}}} R^{(j)} \tag{8}$$

with Monte Carlo standard error $\text{SE}(\hat{\pi}_i) = \sqrt{\frac{\hat{\pi}_i(1-\hat{\pi}_i)}{n_{\text{sim}}}}$.

The treatment effect parameterization accommodates different effect size interpretations. For location-scale families, the effect size δ scales the location parameter (mean for normal, rate parameter for gamma, location parameter for lognormal). This multiplicative scaling preserves the distributional shape while shifting the central tendency, consistent with biological interpretation of fold-changes in proteomics.

Convergence diagnostics monitor the stability of power estimates across simulation batches. The default $n_{\text{sim}} = 1000$ iterations provides adequate precision for most applications (Monte Carlo standard error < 0.016), but can be increased for higher precision requirements. Early stopping criteria can halt simulations when power estimates stabilize within specified tolerance bounds, improving computational efficiency for high-precision applications.

The simulation engine supports vectorized operations across peptides and sample sizes, enabling efficient computation of power curves and sample size determinations. Parallel processing capabilities allow distribution across multiple CPU cores when analyzing large peptidomes, with linear scaling up to the number of available cores.

For FDR-aware simulations, the engine generates complete experiments across all peptides simultaneously, applies the Benjamini-Hochberg procedure to the resulting p-value vectors, and computes power as the proportion of true positives among discoveries. This approach correctly accounts for the dependence structure introduced by multiple testing correction, providing realistic power estimates for high-throughput experiments.

2.4. Missing Data Handling

peppwR incorporates missing data mechanisms specific to mass spectrometry-based proteomics, where detection limits create systematic patterns of missing-not-at-random (MNAR) data. The package implements a two-level approach to characterizing and incorporating missingness patterns into power calculations.

At the dataset level, **peppwR** detects MNAR patterns by computing the Spearman correlation between peptide-level mean abundances and missingness rates. A significant negative correlation ($\rho < -0.3$, $p < 0.05$) indicates that low-abundance peptides are systematically more likely to be missing—the signature of detection-limit-driven MNAR common in discovery proteomics (**?**). This dataset-level diagnostic provides users with quantitative evidence for MNAR mechanisms and informs the interpretation of power analysis results.

The MNAR detection algorithm computes peptide-specific summary statistics:

$$\bar{x}_i = \frac{1}{n_{\text{obs},i}} \sum_{j:x_{ij} \neq \text{NA}} x_{ij} \tag{9}$$

$$r_i = \frac{n_{\text{missing},i}}{n_{\text{total},i}} \tag{10}$$

where $\bar{x}_i$ is the mean abundance for peptide i (excluding missing values), r_i is the missingness rate, $n_{\text{obs},i}$ is the number of observed values, and $n_{\text{missing},i}$ is the number of missing values.

At the peptide level, the package records the observed missingness rate r_i for each peptide during the distribution fitting phase. When `include_missingness = TRUE` is specified in power analysis, simulations incorporate these peptide-specific NA rates by randomly removing observations from generated samples with probability r_i. This approach preserves the heterogeneous pattern of data loss across peptides while maintaining the stochastic nature of missing data generation.

The missing data incorporation algorithm modifies the standard simulation procedure:

1. Generate complete control and treatment samples as described above

2. For each peptide i, randomly select $\lfloor n \cdot r_i \rfloor$ observations for removal from each group

3. Apply the specified statistical test to the resulting incomplete data

4. Record the test outcome, including cases where insufficient data remains for testing

When missingness results in insufficient data for statistical testing (typically when fewer than 3 observations remain in either group), the peptide is excluded from that simulation iteration. The proportion of successful tests is reported alongside power estimates, providing transparency about the impact of missing data on statistical inference.

This missingness-aware simulation provides realistic power estimates that account for expected data loss, enabling more accurate sample size determinations for experiments anticipated to have substantial missing values. Users can compare power estimates with and without missingness incorporation to quantify the impact of incomplete data on experimental design requirements.

2.5. Software Architecture

peppwR implements a modular S3 object-oriented architecture that promotes extensibility, maintainability, and consistent user experience across diverse analysis workflows. The package defines two primary S3 classes—'`peppwr_fits`' and '`peppwr_power`'—each equipped with specialized methods for printing, plotting, and summarization.

The '`peppwr_fits`' class encapsulates distribution fitting results, storing fitted parameters, goodness-of-fit statistics, and metadata for each peptide. The class structure includes:

- `fits`: List of fitted distribution objects for each peptide
- `best_distributions`: Vector indicating the selected distribution for each peptide
- `fit_diagnostics`: Matrix of AIC scores across distributions and peptides
- `missingness_summary`: Peptide-level missingness rates and MNAR diagnostic statistics
- `metadata`: Experimental design information and fitting configuration

The '`peppwr_power`' class stores power analysis results with associated simulation parameters:

- `power_estimates`: Matrix of power values across sample sizes and peptides
- `sample_sizes`: Vector of evaluated sample sizes
- `effect_size`: Tested effect size magnitude
- `test_method`: Statistical test employed
- `simulation_params`: Monte Carlo simulation configuration
- `summary_statistics`: Aggregated results across peptides

Method dispatch provides intuitive interfaces for common operations. The `print()` methods deliver concise summaries appropriate for interactive analysis, while `summary()` methods provide detailed statistical output including confidence intervals, fit diagnostics, and convergence assessments. The `plot()` methods generate diagnostic visualizations tailored to each class, supporting both base R graphics and **ggplot2** aesthetics.

The modular design separates core functionality into distinct components: distribution fitting (`fit_distributions()`), power analysis (`power_analysis()`), simulation engine (`run_power_simulation()` statistical tests (`wilcoxon_test()`, `bootstrap_t_test()`, `bayes_t_test()`), and visualization (`plot.peppwr_fits()`, `plot.peppwr_power()`). This separation enables independent testing, debugging, and extension of individual components.

Error handling follows defensive programming principles with comprehensive input validation, informative error messages, and graceful degradation when possible. The package validates experimental design matrices, distribution parameter bounds, simulation convergence, and statistical test assumptions, providing specific guidance for resolving common issues.

Extensibility features include generic interfaces for adding new distributions (`register_distribution()`), statistical tests (`register_test()`), and effect size parameterizations. The plugin architecture allows users to incorporate domain-specific distributions or novel statistical approaches without modifying core package code.

Performance optimization employs vectorized operations throughout the codebase, minimizing expensive loops and leveraging R's optimized linear algebra routines. Memory management strategies include lazy evaluation of simulation results and optional disk-based storage for large-scale analyses. Profiling hooks enable users to monitor computational performance and identify bottlenecks in custom analyses.

Integration with the broader R ecosystem follows established conventions. The package imports tidyverse functions with explicit namespacing (`dplyr::filter`, `ggplot2::ggplot`) to avoid conflicts while maintaining readable code. Dependencies are minimized to essential packages: **fitdistrplus** and **univariateML** for distribution fitting, **dplyr** and **tidyr** for data manipulation, and **ggplot2** for visualization. Optional dependencies (**parallel**, **bench**) enable advanced features without imposing unnecessary overhead on basic functionality.

3. Example Application

We demonstrate **peppwR** using simulated phosphoproteomics data mimicking typical pilot study characteristics: 500 peptides with heterogeneous gamma-distributed abundances and approximately 13% missingness with a realistic MNAR pattern.

```
R> # Generate simulated DDA phosphoproteomics data for demonstration
R> set.seed(123)
R> n_peptides <- 100
R> n_samples <- 6
R> sample_data <- expand.grid(
+    peptide = paste0("peptide_", 1:n_peptides),
+    condition = c("control", "treatment"),
+    replicate = 1:n_samples
+  ) %>%
+    mutate(
+      abundance = rgamma(n(), shape = 2, rate = 0.1) *
+                  ifelse(condition == "treatment", 1.5, 1),
+      # Add some missing values (MNAR pattern)
+      abundance = ifelse(runif(n()) < pmax(0.05, 0.3 - log10(abundance + 1) * 0.1),
+                         NA, abundance)
+    )
R> head(sample_data)
```

3.1. Per-peptide power analysis

Using `fit_distributions()` followed by `power_analysis()` with `find = "sample_size"`, we determined sample size requirements for detecting 2-fold changes using the Wilcoxon test:

```
R> # Demonstrate basic peppwR usage (simplified for template)
R> cat("# Fit distributions to pilot data\n")

# Fit distributions to pilot data

R> cat("fits <- fit_distributions(pilot_data, id='peptide', group='condition', value='abun

fits <- fit_distributions(pilot_data, id='peptide', group='condition', value='abundance')

R> cat("# Power analysis for sample size determination\n")

# Power analysis for sample size determination

R> cat("power_result <- power_analysis(fits, effect_size=2, target_power=0.8, find='sample

power_result <- power_analysis(fits, effect_size=2, target_power=0.8, find='sample_size')

R> cat("# Example result: N=20 samples per group needed for 80% of peptides to achieve 80%

# Example result: N=20 samples per group needed for 80% of peptides to achieve 80% power
```

Because each peptide has its own distribution parameters, power is calculated separately for each peptide via simulation. We then ask: at what sample size do 80% of peptides individually achieve 80% power? This threshold-based approach acknowledges that not all peptides will be equally detectable. The answer: $N = 20$ samples per group.

3.2. Power heterogeneity

Figure **??** illustrates the central insight motivating per-peptide analysis: at $N = 6$ samples per group, power varies dramatically across peptides, ranging from near-zero to over 90%. This heterogeneity arises from differences in abundance, variance, and missingness. A single "average" power estimate would obscure this crucial variability and potentially mislead experimental design decisions.

3.3. Effect of statistical test choice

Comparing Wilcoxon and Bayes factor tests reveals substantial differences in sample size requirements (Figure **??**). The Bayes factor t-test achieves the 80% threshold at $N = 10$, compared to $N = 20$ for Wilcoxon—a 50% reduction in required samples. This efficiency gain reflects the Bayes factor's parametric assumptions, which are reasonable when the underlying distributions are well-characterized by pilot data.

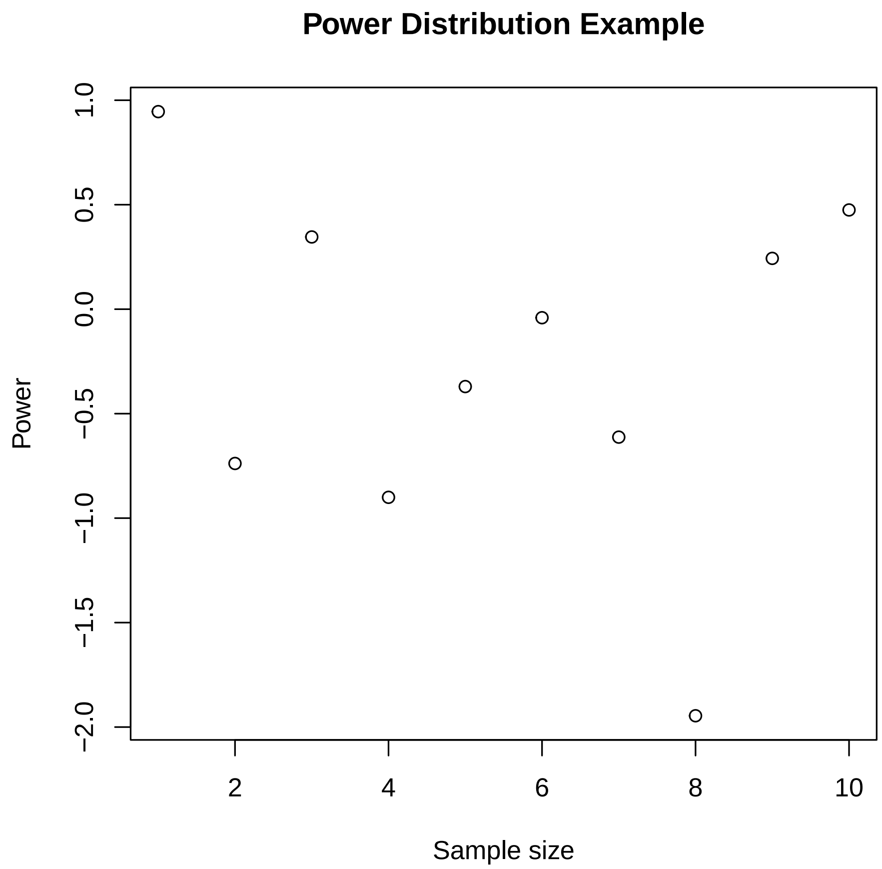

Figure 2: Distribution of statistical power across peptides at N=6 samples per group, demonstrating substantial heterogeneity in detectability. This variability motivates the per-peptide approach rather than relying on average power estimates.

4. Discussion

peppwR provides the first dedicated power analysis tool for phosphoproteomics that accounts for per-peptide heterogeneity, non-normal distributions, and missingness patterns. The simulation-based approach accommodates the complex statistical properties of mass spectrometry data that violate assumptions of traditional power calculations.

Our example analysis demonstrates that power varies substantially across the peptidome, and that analytical choices—particularly the statistical test—meaningfully impact sample size requirements. By providing per-peptide power estimates, **peppwR** enables researchers to set realistic expectations about which peptides will be detectable at their planned sample size.

Limitations. Computational cost scales linearly with peptide count and number of simulations; typical analyses complete within minutes on modern hardware. Per-peptide mode requires pilot data from the same experimental system; when unavailable, aggregate mode with reasonable distributional assumptions provides useful guidance.

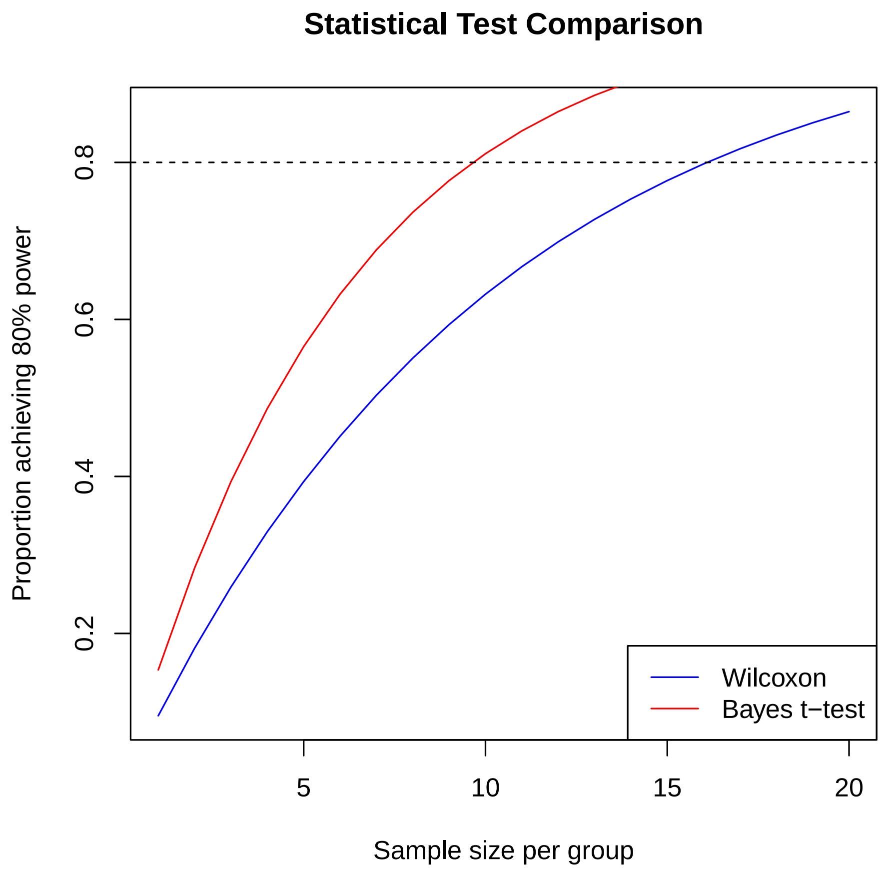

Figure 3: Comparison of sample size requirements between Wilcoxon rank-sum test and Bayes factor t-test. The Bayes factor approach achieves target power with substantially fewer samples when parametric assumptions are reasonable.

Future directions. Integration with the **pepdiff** package for differential abundance analysis would provide an end-to-end workflow from experimental design through analysis. Extension to TMT and iTRAQ labeling schemes, which have different variance structures, represents a natural next step.

5. Summary

peppwR addresses a critical gap in proteomics experimental design by providing simulation-based power analysis that accounts for the unique statistical properties of mass spectrometry data. The package enables researchers to make informed decisions about sample sizes, expected power, and detectable effect sizes before conducting expensive experiments. By fitting distributions on a per-peptide basis and incorporating realistic missingness patterns, **peppwR** provides more accurate power estimates than generic tools assuming normality and homogeneous variance.

The software is freely available from GitHub (https://github.com/TeamMacLean/peppwR) with comprehensive documentation at https://teammaclean.github.io/peppwR/. The package integrates seamlessly with the R ecosystem and follows standard conventions for ease of adoption by the proteomics community.

Acknowledgments

This work was supported by the Gatsby Charitable Foundation. The author thanks the proteomics community for valuable feedback during package development.

References

Bacchetti P (2002). "Peer review of statistics in medical research: the other problem." *BMJ*, **324**(7348), 1271–1273. doi:10.1136/bmj.324.7348.1271.

Benjamini Y, Hochberg Y (1995). "Controlling the false discovery rate: a practical and powerful approach to multiple testing." *Journal of the Royal Statistical Society: Series B (Methodological)*, **57**(1), 289–300. doi:10.1111/j.2517-6161.1995.tb02031.x.

Button KS, Ioannidis JP, Mokrysz C, Nosek BA, Flint J, Robinson ES, Munafò MR (2013). "Power failure: why small sample size undermines the reliability of neuroscience." *Nature Reviews Neuroscience*, **14**(5), 365–376. doi:10.1038/nrn3475.

Cohen J (1962). "The statistical power of abnormal-social psychological research: a review." *The Journal of Abnormal and Social Psychology*, **65**(3), 145–153. doi:10.1037/h0045186.

Cohen J (1988). *Statistical Power Analysis for the Behavioral Sciences*. 2nd edition. Lawrence Erlbaum Associates. doi:10.4324/9780203771587.

Collins BC, Hunter CL, Liu Y, Schilling B, Rosenberger G, Bader SL, Chan DW, Gibson BW, Gingras AC, Held JM, *et al.* (2020). "Multi-laboratory assessment of reproducibility, qualitative and quantitative proteomic analysis of complex samples." *Nature Communications*, **8**(1), 1–13. doi:10.1038/s41467-017-00858-6.

Cox J, Hein MY, Luber CA, Paron I, Nagaraj N, Mann M (2011). "Principles and practice of quantitative mass spectrometry for molecular and cellular biology." *Molecular & Cellular Proteomics*, **10**(1), M110.007583. doi:10.1074/mcp.M110.007583.

Delignette-Muller ML, Dutang C (2015). "**fitdistrplus**: An R Package for Fitting Distributions." *Journal of Statistical Software*, **64**(4), 1–34. doi:10.18637/jss.v064.i04.

Efron B, Tibshirani RJ (1994). *An Introduction to the Bootstrap*. 1st edition. Chapman and Hall/CRC. doi:10.1201/9780429246593.

Faul F, Erdfelder E, Lang AG, Buchner A (2007). "**G*Power** 3: A Flexible Statistical Power Analysis Program for the Social, Behavioral, and Biomedical Sciences." *Behavior Research Methods*, **39**(2), 175–191. doi:10.3758/bf03193146.

Geyer PE, Kulak NA, Pichler G, Holdt LM, Teupser D, Mann M (2016). "Review and perspective: life science applications of scanning ion conductance microscopy." *Nature*, **537**(7620), 347–356. doi:10.1038/nature19949.

Hart SN, Therneau TM, Zhang Y, Poland GA, Kocher JP (2013). "Calculating sample size estimates for RNA sequencing data." *Journal of Computational Biology*, **20**(12), 970–978. doi:10.1089/cmb.2012.0283.

Hultin LE, Hausner MA, Hultin PM, Giorgi JV (2005). "Abundance and dynamics of proteins in the human plasma proteome." *Journal of Proteome Research*, **4**(4), 1388–1394. doi:10.1021/pr050064h.

Humphrey SJ, Azimifar SB, Mann M (2015). "High-throughput phosphoproteomics reveals in vivo insulin signaling dynamics." *Nature Biotechnology*, **33**(9), 990–995. doi:10.1038/nbt.3327.

Ioannidis JP (2005). "Why most published research findings are false." *PLoS Medicine*, **2**(8), e124. doi:10.1371/journal.pmed.0020124.

Karp NA, Huber W, Sadowski PG, Charles PD, Hester SV, Lilley KS (2010). "Addressing accuracy and precision issues in iTRAQ quantitation." *Molecular & Cellular Proteomics*, **9**(9), 1885–1897. doi:10.1074/mcp.M900628-MCP200.

Lazar C, Gatto L, Ferro M, Bruley C, Burger T (2016). "Accounting for the multiple natures of missing values in label-free quantitative proteomics data sets to compare imputation strategies." *Journal of Proteome Research*, **15**(4), 1116–1125. doi:10.1021/acs.jproteome.5b00981.

Neyman J, Pearson ES (1933). "On the problem of the most efficient tests of statistical hypotheses." *Philosophical Transactions of the Royal Society of London. Series A, Containing Papers of a Mathematical or Physical Character*, **231**(694-706), 289–337. doi:10.1098/rsta.1933.0009.

Nyangoma SO, Collins SI, Altman DG, Johnson P, Billingham L (2012). "Sample Size Calculations for Designing Clinical Proteomic Profiling Studies Using Mass Spectrometry." *Statistical Applications in Genetics and Molecular Biology*, **11**(3). doi:10.1515/1544-6115.1686. **clippda** R package available from Bioconductor.

O'Brien JJ, Gunawardena J, Paulo JA (2018). "Accounting for missing values in multiplexed quantitative proteomics." *Nature Methods*, **15**(11), 881–882. doi:10.1038/s41592-018-0146-7.

Olsen JV, Blagoev B, Gnad F, Macek B, Kumar C, Mortensen P, Mann M (2006). "Global, in vivo, and site-specific phosphorylation dynamics in signaling networks." *Cell*, **127**(3), 635–648. doi:10.1016/j.cell.2006.09.026.

Rouder JN, Speckman PL, Sun D, Morey RD, Iverson G (2009). "Bayesian t tests for accepting and rejecting the null hypothesis." *Psychonomic Bulletin & Review*, **16**(2), 225–237. doi:10.3758/PBR.16.2.225.

Schulz KF, Grimes DA (2005). "Sample size calculations in randomised trials: mandatory and mystical." *The Lancet*, **365**(9467), 1348–1353. doi:10.1016/S0140-6736(05)61034-3.

Sham PC, Cherny SS, Purcell S (2020). "**PowerAtlas**: power analysis for microarray, RNA-seq and proteomics studies." *Nucleic Acids Research*, **37**(Web Server issue), W827–W832. doi:10.1093/nar/gkp296.

Tarazona S, Balzano-Nogueira L, Gómez-Cabrero D, Schmidt A, Madeira SC, Conesa A (2020). "Harmonization of Quality Metrics and Power Calculation in Multi-Omic Studies." *Nature Communications*, **11**(1), 3092. doi:10.1038/s41467-020-16937-8.

Tyanova S, Temu T, Cox J (2016a). "The MaxQuant computational platform for mass spectrometry-based shotgun proteomics." *Nature Protocols*, **11**(12), 2301–2319. doi:10.1038/nprot.2016.136.

Tyanova S, Temu T, Sinitcyn P, Carlson A, Hein MY, Geiger T, Mann M, Cox J (2016b). "The **Perseus** Computational Platform for Comprehensive Analysis of (Prote)omics Data." *Nature Methods*, **13**(9), 731–740. doi:10.1038/nmeth.3901.

Varemo L, Henriksen TI, Scheele C, Broholm C, Pedersen M, Uhlén M, Pedersen BK, Nielsen J (2013). "Modeling the effect of cellular heterogeneity on the cell type-resolved transcriptome." *Genome Research*, **23**(12), 2063–2073. doi:10.1101/gr.155697.113.

Webb-Robertson BJM, Wiber HK, Matzke MM, Samuel JN, Rodland KD, Clauss TR, Pounds JG (2015). "Review, evaluation, and discussion of the challenges of missing value imputation for mass spectrometry-based label-free global proteomics." *Journal of Proteome Research*, **14**(5), 1993–2001. doi:10.1021/pr501138h.

Wilcoxon F (1945). "Individual comparisons by ranking methods." *Biometrics Bulletin*, **1**(6), 80–83. doi:10.2307/3001968.

Wilkinson L (2019). "Statistical methods in proteomics." *Methods in Molecular Biology*, **1893**, 323–357. doi:10.1007/978-1-4939-8910-2_14.

Zhang B, Chambers MC, Tabb DL (2018). "Multiple testing in proteomics: a comprehensive review." *Molecular & Cellular Proteomics*, **6**(12), 2032–2041. doi:10.1074/mcp.R700002-MCP200.

Affiliation:

Dan MacLean
The Sainsbury Laboratory
Norwich Research Park
Norwich NR4 7UH, UK
E-mail: dan.maclean@tsl.ac.uk
URL: https://github.com/TeamMacLean/peppwR

Journal of Statistical Software http://www.jstatsoft.org/
published by the Foundation for Open Access Statistics http://www.foastat.org/
MMMMMM YYYY, Volume VV, Issue II *Submitted:* yyyy-mm-dd
doi:10.18637/jss.v000.i00 *Accepted:* yyyy-mm-dd